Introduction to R for SAS programmers

Sadchla Mascary, Stefan Thoma, Zelos Zhu, Thomas Neitman

1/17/23

Table of contents

About

On this page you find the materials for the workshop *Introduction to R for SAS programmers*. The workshop provides an RStudio cloud workspace, so you don't need to have R and RStudio locally installed.

If you want to participate in this workshop using your local machine, you need to download the data first. Executing the following chunk in R on your machine will create and populate a data folder in your current working directory. If you work on the RStudio cloud instance provided by us via link, you do not need to run this chunk.

```
# set paths and data names
external.path <-
  "https://github.com/pharmaverse/intro-to-r-for-sas-programmers-workshop/blob/main/data"

local.path <- ("data")
subdir <- file.path(local.path, "save_data")
files <- c(
  "adsl.RData",
  "adsl.csv",
  "adsl.sas7bdat",
  "adsl_1.RData",
  "adsl_2.RData",
  "ae.rds",
  "dm.rds",
  "ds.rds",
  "ex.rds",
  "suppdm.rds",
  "suppds.rds"
  )

# external files (with path)
urls <- file.path(external.path, paste(files, "raw=true", sep = "?"))
# local files (with path)
dest <- file.path(local.path, files)

# create data folder in wd
```

```r
if (!file.exists(local.path)) {
  dir.create(local.path)
  # subdirectory
  dir.create(subdir)
}

# download files if needed
download.file(urls, destfile = dest)
```

Further material

You can't get enough? Here is a resource to help transitioning from SAS to R.

Part I

datatype & structure

1 datatype and structure

Sadchla Mascary

R can be used as a calculator following the order of operations using the basic arithmetic operators, although, the arithmetic equal sign (=) in the equivalent of ==.

```r
# simple calculations
3*2
```

```
[1] 6
```

```r
(59 + 73 + 2) / 3
```

```
[1] 44.66667
```

```r
# complex calculations
pi/8
```

```
[1] 0.3926991
```

1.1 Storing outputs

An object can be created to assign the value of your operation to a specific variable name, which can be reused later in the R session. Using the `object_name <- value` naming convention, you can assign (`<-`) the value (`(59 + 73 + 2) / 3`) to an `object_name` simple_cal to look like `simple_cal <- (59 + 73 + 2) / 3` to store the evaluation of that calculation.

```r
x <- 1:10
```

```r
y <- 2*x

simple_cal <- (59 + 73 + 2) / 3
```

1.2 Loading data into R

Depending on the formats for the files containing your data, we can use different base R functions to read and load data into memory

R has two native data formats, **Rdata** (sometimes call Rda) and **RDS**.

Rdata can be selected R objects or a workspace, and **RDS** are single R object. R has base functions available to read the two native data formats, and some delimited files.

```r
# saving rdata
save(x, file = "data/intro_1.RData")
# Save multiple objects
save(x, y, file = "data/intro_2.RData")

# Saving the entire workspake
save.image(file="data/intro_program.RData")

# We can follow the syntax for saving single Rdata object to save Rds files
# saveRDS(object, file = "my_data.rds")


# loading Rdata or Rda files
load(file = "data/intro_program.RData")

# loading RDS
# We can follow the syntax for read Rdata object to sread Rds files using the readRDS()


# Comma delimited
adsl_CSV <- read.csv("data/adsl.csv", header = TRUE)

# Save CSV
adsl_csv_save <- write.csv(adsl_CSV, "data/save_data/adsl.csv", row.names=TRUE)


adsl_TAB_save <- write.table(
```

```
    adsl_CSV,
    "data/save_data/adsl.txt",
    append = FALSE,
    sep = "\t",
    dec = ".",
    row.names = TRUE,
    col.names = TRUE
)


# Tab-delimited
adsl_TAB <- read.table("data/save_data/adsl.txt", header = TRUE, sep = "\t")
```

1.3 R Packages

R packages are a collection of reusable functions, compiled codes, documentation, sample data and tests. Some formats of data require the use of an R package in order to load that data into memory. Share-able R packages are typically stored in a repository such as the Comprehensive R Archive Network (CRAN), Bioconductor, and GitHub.

1.4 Installing R packages

```
# From CRAN
#install.packages("insert_package_name")
# {haven} is used to import or export foreign statistical format files (SPSS, Stata, SAS)
install.packages("haven")

# {readxl}
install.packages("readxl")


# From Github
remotes::install_github("pharmaverse/admiral", ref = "devel")
```

1.5 Using R packages, functions from an R package, and accessing help pages

Since an R packages are a collection of functions, you can choose to load the entire package within R memory or just the needed function from that package. Usually, the order you choose to load your package does not make a difference, unless you are loading two or more packages that has functions with the same name. If you are loading two or more packages with common function name, then the package loaded last will hide that function in the earlier packages, so in that case is important to note the order you choose to load the packages.

```r
# Read file using ::
adsl_sas1 <- haven::read_sas("data/adsl.sas7bdat")

# read file using library call
library(haven)
adsl_sas2 <- read_sas("data/adsl.sas7bdat")


# Reading Excel xls|xlsx files
# read_excel reads both xls|xlsx files but read_xls and read_xlsx can also be used to read

# if NA are represented by another something other than blank then you can specified the N
# within the read_excel() function
```

1.6 Data types

R has different types of **Datatype**
* Integer * numeric * Character * Logical * complex * raw

But we will focus on the top 4.

```r
set.seed(1234)

type_int <- (1:5)
type_num <- rnorm(5)
type_char <- "USUBJID"
type_logl_1 <- TRUE
type_logl_2 <- FALSE
```

```r
class(type_int)
```

```
[1] "integer"
```

```r
class(type_num)
```

```
[1] "numeric"
```

```r
class(type_logl_1)
```

```
[1] "logical"
```

```r
class(type_logl_2)
```

```
[1] "logical"
```

```r
class(type_char)
```

```
[1] "character"
```

1.7 Date formats

There are base R functions that can be used to format a date object similar to the Date9 formatted date variable from SAS. In addition, there are R packages available, such as {lubridate}, for more complex date/date time formatted objects.

```r
# using adsl_sas1 RFSTDTC
class(adsl_sas1$RFSTDTC)
```

```
[1] "character"
```

```r
# Convert the date from that adsl_sas1 into a date variable
adsl_sas1$RFSTDTC <- as.Date(adsl_sas1$RFSTDTC)
class(adsl_sas1$RFSTDTC)
```

```
[1] "Date"
```

```r
date9 <- lubridate::as_date(18757)
lubridate::mdy(adsl_sas1$RFSTDTC)
```

```
Warning: All formats failed to parse. No formats found.
```

```
  [1] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
 [26] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
 [51] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
 [76] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
[101] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
[126] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
[151] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
[176] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
[201] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
[226] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
[251] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
[276] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
[301] NA NA NA NA NA NA
```

1.8 Structures

Data structures are dimensional ways of organizing the data. There are different data structures in R, let's focus on **vectors** and **dataframe**

Vectors are 1 dimensional collection of data that can contain one or more element of the same data type

```r
vect_1 <- 2
vect_2 <- c(2, "USUBJID")

class(vect_1)
```

```
[1] "numeric"
```

```
class(vect_2)
```

[1] "character"

```
# Saving vectors from a dataset to a specific variable
usubjid <- adsl_sas1$USUBJID
subjid <- adsl_sas1[, 3]
```

Dataframe is similar to SAS data sets and are 2 dimensional collection of vectors. Dataframe can store vectors of different types but must be of the same length

```
df <- data.frame(
  age = c(65, 20, 37,19,45),
  seq = (1:5),
  type_logl = c(TRUE,FALSE, TRUE, TRUE, FALSE),
  usubjid = c("001-940-9785","002-950-9726","003-940-9767","004-940-9795","005-940-9734")
)

# str() provides the data structure for each object in the dataframe
str(df)
```

```
'data.frame':   5 obs. of  4 variables:
$ age      : num   65 20 37 19 45
$ seq      : int   1 2 3 4 5
$ type_logl: logi  TRUE FALSE TRUE TRUE FALSE
$ usubjid  : chr   "001-940-9785" "002-950-9726" "003-940-9767" "004-940-9795" ...
```

```
# In addition to the data structure per variable, also get some descriptive statistics
summary(df)
```

```
     age               seq      type_logl          usubjid
 Min.   :19.0    Min.   :1    Mode :logical    Length:5
 1st Qu.:20.0    1st Qu.:2    FALSE:2          Class :character
 Median :37.0    Median :3    TRUE :3          Mode  :character
 Mean   :37.2    Mean   :3
 3rd Qu.:45.0    3rd Qu.:4
 Max.   :65.0    Max.   :5
```

2 datatype and structure exercise

Install and load the following packages

{tidyverse} {admiral} {dplyr} {tidyr} {admiral.test}

```
#installing the packages
install.packages(c("tidyverse", "admiral", "dplyr", "tidyr"))

library(tidyverse)
library(admiral)
library(admiral.test)
library(dplyr)
library(tidyr)
```

3 Exercise 2

Import **adsl.sas7bdat** as adsl

```r
adsl <-  haven::read_sas("data/adsl.sas7bdat")
```

Part II

dplyr

4 dplyr

Stefan Thoma

The `tidyverse` is a collection of R packages designed for data science. It includes packages such as `ggplot2` for data visualization, `dplyr` for data manipulation, and `tidyr` for reshaping data. The `tidyverse` is built around the idea of "tidy data," which is a standardized way of organizing and structuring data for analysis. The packages in the `tidyverse` are designed to work together seamlessly, making it a popular choice for data scientists and analysts who use R.

```
library(tidyverse)
```

```
-- Attaching packages --------------------------------------- tidyverse 1.3.2 --
v ggplot2 3.3.6      v purrr   0.3.5
v tibble  3.1.8      v dplyr   1.0.10
v tidyr   1.2.1      v stringr 1.4.1
v readr   2.1.3      v forcats 0.5.2
-- Conflicts ------------------------------------------ tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
```

Read data

```
adsl <- read_csv("data/adsl.csv")
```

```
Rows: 306 Columns: 50
-- Column specification --------------------------------------------------------
Delimiter: ","
chr  (23): STUDYID, USUBJID, DTHFL, AGEU, SEX, RACE, ETHNIC, ARMCD, ARM, ACT...
dbl   (7): SUBJID, SITEID, AGE, DMDY, TRTDURD, DTHADY, LDDTHELD
lgl   (3): RFICDTC, REGION1, DTHA30FL
dttm  (3): RFPENDTC, TRTSDTM, TRTEDTM
```

```
date (14): RFSTDTC, RFENDTC, RFXSTDTC, RFXENDTC, DTHDTC, DMDTC, TRTSDT, TRTE...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

We can have a look at the data using many different commands / functions, e.g. the `head()` function which gives us the first six observations:

```
head(adsl)
```

```
# A tibble: 6 x 50
  STUDYID     USUBJID SUBJID RFSTDTC    RFENDTC    RFXSTDTC   RFXENDTC   RFICDTC
  <chr>       <chr>    <dbl> <date>     <date>     <date>     <date>     <lgl>
1 CDISCPILOT~ 01-701~   1015 2014-01-02 2014-07-02 2014-01-02 2014-07-02 NA
2 CDISCPILOT~ 01-701~   1023 2012-08-05 2012-09-02 2012-08-05 2012-09-01 NA
3 CDISCPILOT~ 01-701~   1028 2013-07-19 2014-01-14 2013-07-19 2014-01-14 NA
4 CDISCPILOT~ 01-701~   1033 2014-03-18 2014-04-14 2014-03-18 2014-03-31 NA
5 CDISCPILOT~ 01-701~   1034 2014-07-01 2014-12-30 2014-07-01 2014-12-30 NA
6 CDISCPILOT~ 01-701~   1047 2013-02-12 2013-03-29 2013-02-12 2013-03-09 NA
# ... with 42 more variables: RFPENDTC <dttm>, DTHDTC <date>, DTHFL <chr>,
#   SITEID <dbl>, AGE <dbl>, AGEU <chr>, SEX <chr>, RACE <chr>, ETHNIC <chr>,
#   ARMCD <chr>, ARM <chr>, ACTARMCD <chr>, ACTARM <chr>, COUNTRY <chr>,
#   DMDTC <date>, DMDY <dbl>, TRT01P <chr>, TRT01A <chr>, TRTSDTM <dttm>,
#   TRTSTMF <chr>, TRTEDTM <dttm>, TRTETMF <chr>, TRTSDT <date>, TRTEDT <date>,
#   TRTDURD <dbl>, SCRFDT <date>, EOSDT <date>, EOSSTT <chr>, FRVDT <date>,
#   RANDDT <date>, DTHDT <date>, DTHADY <dbl>, LDDTHELD <dbl>, ...
```

4.1 `dplyr`

`dplyr` is a package in the `tidyverse` that provides a set of functions for efficiently manipulating and cleaning data. It is built around the idea of "verbs" that correspond to common data manipulation tasks, such as `select()` for selecting specific columns from a data frame, `filter()` for filtering rows based on certain conditions, `arrange()` for sorting data-frames and `group_by()` and `summarize()` for grouping and summarizing data by one or more variables.

`dplyr` is not strictly needed for any of that, everything can be done in base R. However, `dplyr` provides a framework to write readable code and a pipeline to work efficiently.

There are various functions within `dplyr` for datawrangling which follow a consistent structure. The first input of the most used `dplyr` functions is the data-frame. Then follow arguments specifying the behaviour of the function. Compared to the base r syntax we do not have

to write column / variable names in quotation marks; `dplyr` syntax lets us refer to columns within a data-frame without the need to always reference the data-frame of origin.

4.1.1 select

The `select` function lets us select all variables mentioned in the arguments (and drops all other variables). Alternatively, we can selectively drop variables if we place a minus (`-`) in front of the variable name.
We can first have a look at all variable names of the data-frame:

```
names(adsl)
```

```
 [1] "STUDYID"  "USUBJID"  "SUBJID"   "RFSTDTC"  "RFENDTC"  "RFXSTDTC"
 [7] "RFXENDTC" "RFICDTC"  "RFPENDTC" "DTHDTC"   "DTHFL"    "SITEID"
[13] "AGE"      "AGEU"     "SEX"      "RACE"     "ETHNIC"   "ARMCD"
[19] "ARM"      "ACTARMCD" "ACTARM"   "COUNTRY"  "DMDTC"    "DMDY"
[25] "TRT01P"   "TRT01A"   "TRTSDTM"  "TRTSTMF"  "TRTEDTM"  "TRTETMF"
[31] "TRTSDT"   "TRTEDT"   "TRTDURD"  "SCRFDT"   "EOSDT"    "EOSSTT"
[37] "FRVDT"    "RANDDT"   "DTHDT"    "DTHADY"   "LDDTHELD" "LSTALVDT"
[43] "AGEGR1"   "SAFFL"    "RACEGR1"  "REGION1"  "LDDTHGR1" "DTH30FL"
[49] "DTHA30FL" "DTHB30FL"
```

And then select the desired variables:

```
# dplyr::select
select(adsl,
       STUDYID,
       USUBJID,
       ARM,
       AGE,
       SEX,
       RACE)
```

```
# A tibble: 306 x 6
  STUDYID     USUBJID     ARM                   AGE SEX   RACE
  <chr>       <chr>       <chr>               <dbl> <chr> <chr>
1 CDISCPILOT01 01-701-1015 Placebo              63 F     WHITE
2 CDISCPILOT01 01-701-1023 Placebo              64 M     WHITE
3 CDISCPILOT01 01-701-1028 Xanomeline High Dose 71 M     WHITE
4 CDISCPILOT01 01-701-1033 Xanomeline Low Dose  74 M     WHITE
```

```
 5 CDISCPILOT01 01-701-1034 Xanomeline High Dose    77 F     WHITE
 6 CDISCPILOT01 01-701-1047 Placebo                 85 F     WHITE
 7 CDISCPILOT01 01-701-1057 Screen Failure          59 F     WHITE
 8 CDISCPILOT01 01-701-1097 Xanomeline Low Dose     68 M     WHITE
 9 CDISCPILOT01 01-701-1111 Xanomeline Low Dose     81 F     WHITE
10 CDISCPILOT01 01-701-1115 Xanomeline Low Dose     84 M     WHITE
# ... with 296 more rows
```

We end up with a new data-frame including only the selected variables. Note here that we do not save the resulting data-frame at the moment.

There are also some helper functions to use within the `select` function of `dplyr`. `starts_with()` `ends_with()` `num_range()`. They allow us to select multiple columns sharing a naming structure. `num_range()` let's us select consecutively numbered columns, e.g.: `num_range("example", 1:4)` would select the columns named: example1, example2, example3, example4.

We can try out `starts_with()`:

```
  select(adsl,
         USUBJID,
         starts_with("trt"))
```

```
# A tibble: 306 x 10
   USUBJID TRT01P TRT01A TRTSDTM             TRTSTMF TRTEDTM             TRTETMF
   <chr>   <chr>  <chr>  <dttm>              <chr>   <dttm>              <chr>
 1 01-701~ Place~ Place~ 2014-01-02 00:00:00 H       2014-07-02 23:59:59 H
 2 01-701~ Place~ Place~ 2012-08-05 00:00:00 H       2012-09-01 23:59:59 H
 3 01-701~ Xanom~ Xanom~ 2013-07-19 00:00:00 H       2014-01-14 23:59:59 H
 4 01-701~ Xanom~ Xanom~ 2014-03-18 00:00:00 H       2014-03-31 23:59:59 H
 5 01-701~ Xanom~ Xanom~ 2014-07-01 00:00:00 H       2014-12-30 23:59:59 H
 6 01-701~ Place~ Place~ 2013-02-12 00:00:00 H       2013-03-09 23:59:59 H
 7 01-701~ Scree~ Scree~ NA                  <NA>    NA                  <NA>
 8 01-701~ Xanom~ Xanom~ 2014-01-01 00:00:00 H       2014-07-09 23:59:59 H
 9 01-701~ Xanom~ Xanom~ 2012-09-07 00:00:00 H       2012-09-16 23:59:59 H
10 01-701~ Xanom~ Xanom~ 2012-11-30 00:00:00 H       2013-01-23 23:59:59 H
# ... with 296 more rows, and 3 more variables: TRTSDT <date>, TRTEDT <date>,
#   TRTDURD <dbl>
```

And `ends_with()`:

```
# in this df, all variables that contain dates end with "DT".
# We can select them:
select(adsl,
       USUBJID,
       ends_with("DT"))
```

```
# A tibble: 306 x 9
   USUBJID     TRTSDT     TRTEDT     SCRFDT     EOSDT      FRVDT      RANDDT
   <chr>       <date>     <date>     <date>     <date>     <date>     <date>
 1 01-701-1015 2014-01-02 2014-07-02 NA         2014-07-02 NA         2014-01-02
 2 01-701-1023 2012-08-05 2012-09-01 NA         2012-09-02 2013-02-18 2012-08-05
 3 01-701-1028 2013-07-19 2014-01-14 NA         2014-01-14 NA         2013-07-19
 4 01-701-1033 2014-03-18 2014-03-31 NA         2014-04-14 2014-09-15 2014-03-18
 5 01-701-1034 2014-07-01 2014-12-30 NA         2014-12-30 NA         2014-07-01
 6 01-701-1047 2013-02-12 2013-03-09 NA         2013-03-29 2013-07-28 2013-02-12
 7 01-701-1057 NA         NA         2013-12-20 NA         NA         NA
 8 01-701-1097 2014-01-01 2014-07-09 NA         2014-07-09 NA         2014-01-01
 9 01-701-1111 2012-09-07 2012-09-16 NA         2012-09-17 2013-02-22 2012-09-07
10 01-701-1115 2012-11-30 2013-01-23 NA         2013-01-23 2013-05-20 2012-11-30
# ... with 296 more rows, and 2 more variables: DTHDT <date>, LSTALVDT <date>
```

If we want a data-frame that does not include any dates, we can make use of the minus sign in combination with the `ends_with()` function:

```
select(adsl,
       -ends_with("DT"))
```

```
# A tibble: 306 x 42
   STUDYID    USUBJID  SUBJID RFSTDTC    RFENDTC    RFXSTDTC   RFXENDTC   RFICDTC
   <chr>      <chr>    <dbl>  <date>     <date>     <date>     <date>     <lgl>
 1 CDISCPILO~ 01-701~   1015  2014-01-02 2014-07-02 2014-01-02 2014-07-02 NA
 2 CDISCPILO~ 01-701~   1023  2012-08-05 2012-09-02 2012-08-05 2012-09-01 NA
 3 CDISCPILO~ 01-701~   1028  2013-07-19 2014-01-14 2013-07-19 2014-01-14 NA
 4 CDISCPILO~ 01-701~   1033  2014-03-18 2014-04-14 2014-03-18 2014-03-31 NA
 5 CDISCPILO~ 01-701~   1034  2014-07-01 2014-12-30 2014-07-01 2014-12-30 NA
 6 CDISCPILO~ 01-701~   1047  2013-02-12 2013-03-29 2013-02-12 2013-03-09 NA
 7 CDISCPILO~ 01-701~   1057  NA         NA         NA         NA         NA
 8 CDISCPILO~ 01-701~   1097  2014-01-01 2014-07-09 2014-01-01 2014-07-09 NA
 9 CDISCPILO~ 01-701~   1111  2012-09-07 2012-09-17 2012-09-07 2012-09-16 NA
10 CDISCPILO~ 01-701~   1115  2012-11-30 2013-01-23 2012-11-30 2013-01-23 NA
# ... with 296 more rows, and 34 more variables: RFPENDTC <dttm>,
```

```
#   DTHDTC <date>, DTHFL <chr>, SITEID <dbl>, AGE <dbl>, AGEU <chr>, SEX <chr>,
#   RACE <chr>, ETHNIC <chr>, ARMCD <chr>, ARM <chr>, ACTARMCD <chr>,
#   ACTARM <chr>, COUNTRY <chr>, DMDTC <date>, DMDY <dbl>, TRT01P <chr>,
#   TRT01A <chr>, TRTSDTM <dttm>, TRTSTMF <chr>, TRTEDTM <dttm>, TRTETMF <chr>,
#   TRTDURD <dbl>, EOSSTT <chr>, DTHADY <dbl>, LDDTHELD <dbl>, AGEGR1 <chr>,
#   SAFFL <chr>, RACEGR1 <chr>, REGION1 <lgl>, LDDTHGR1 <chr>, ...
```

> **i** We can use the `select()` function to reorder the variables in the data-frame. This
> does not affect the order of rows.
>
> ```
> select(adsl,
>        ARM,
>        USUBJID)
> ```
>
> ```
> # A tibble: 306 x 2
>    ARM                 USUBJID
>    <chr>               <chr>
>  1 Placebo             01-701-1015
>  2 Placebo             01-701-1023
>  3 Xanomeline High Dose 01-701-1028
>  4 Xanomeline Low Dose  01-701-1033
>  5 Xanomeline High Dose 01-701-1034
>  6 Placebo             01-701-1047
>  7 Screen Failure      01-701-1057
>  8 Xanomeline Low Dose  01-701-1097
>  9 Xanomeline Low Dose  01-701-1111
> 10 Xanomeline Low Dose  01-701-1115
> # ... with 296 more rows
> ```

4.1.2 filter

The `filter` function allows us to look at a subset of observations. As input, the function
requires a logical vector and (of course) a data-frame. This time, we first save the reduced
(selected) data-frame and use that as the first argument to `filter`.

```
selected_data <- select(adsl,
                        STUDYID,
                        USUBJID,
                        ARM,
                        AGE,
```

```
                        SEX,
                        RACE)
```

The logical vector is generally created within the function call and can use any of the following logic operators:

```
<               less than
<=              less than or equal to
>               greater than
>=              greater than or equal to
==              equal
!=              not equal
!x              not x (negation)
x | y           x OR y
x & y           x AND y
x %in% y        logical vector of length x with TRUE if element of x is in y
```

Within `filter`, we can chain logical vectors by separating them with a comma (`,`). Lets have a look at women that are 70 and older:

```r
filter(selected_data,
       AGE >= 70,
       SEX == "F")
```

```
# A tibble: 141 x 6
   STUDYID      USUBJID     ARM                  AGE SEX   RACE
   <chr>        <chr>       <chr>              <dbl> <chr> <chr>
 1 CDISCPILOT01 01-701-1034 Xanomeline High Dose  77 F     WHITE
 2 CDISCPILOT01 01-701-1047 Placebo               85 F     WHITE
 3 CDISCPILOT01 01-701-1111 Xanomeline Low Dose   81 F     WHITE
 4 CDISCPILOT01 01-701-1133 Xanomeline High Dose  81 F     WHITE
 5 CDISCPILOT01 01-701-1146 Xanomeline High Dose  75 F     WHITE
 6 CDISCPILOT01 01-701-1153 Placebo               79 F     WHITE
 7 CDISCPILOT01 01-701-1162 Screen Failure        82 F     WHITE
 8 CDISCPILOT01 01-701-1181 Xanomeline High Dose  79 F     WHITE
 9 CDISCPILOT01 01-701-1192 Xanomeline Low Dose   80 F     WHITE
10 CDISCPILOT01 01-701-1203 Placebo               81 F     BLACK OR AFRICAN A~
# ... with 131 more rows
```

Now we have a reduced data frame with female patients over 70. However, the nested call is not very intuitive to read. If any more functions get added to this code, it becomes even less readable. That is where the pipe operator (`%>%`) comes in.

> **i** The pipe operator let us chain multiple `dplyr` commands, so we can always forward the previously filtered / selected / arranged dataframe and keep working with it. The pipe operator let's us write nested function calls in a sequential way. Traditionally, we start a new line after every pipe operator.

```r
# select, filter, & pipe:
adsl %>% # This pipe forwards adsl to the select function as its first argument
  select(STUDYID,
         USUBJID,
         ARM,
         AGE,
         SEX,
         RACE) %>% # this pipe forwards the selected variables to the filter function
  filter(AGE >= 70,
         SEX == "F")
```

```
# A tibble: 141 x 6
   STUDYID     USUBJID     ARM                    AGE SEX   RACE
   <chr>       <chr>       <chr>                <dbl> <chr> <chr>
 1 CDISCPILOT01 01-701-1034 Xanomeline High Dose    77 F     WHITE
 2 CDISCPILOT01 01-701-1047 Placebo                 85 F     WHITE
 3 CDISCPILOT01 01-701-1111 Xanomeline Low Dose     81 F     WHITE
 4 CDISCPILOT01 01-701-1133 Xanomeline High Dose    81 F     WHITE
 5 CDISCPILOT01 01-701-1146 Xanomeline High Dose    75 F     WHITE
 6 CDISCPILOT01 01-701-1153 Placebo                 79 F     WHITE
 7 CDISCPILOT01 01-701-1162 Screen Failure          82 F     WHITE
 8 CDISCPILOT01 01-701-1181 Xanomeline High Dose    79 F     WHITE
 9 CDISCPILOT01 01-701-1192 Xanomeline Low Dose     80 F     WHITE
10 CDISCPILOT01 01-701-1203 Placebo                 81 F     BLACK OR AFRICAN A~
# ... with 131 more rows
```

There is another inline operator which can be very useful within the filter function; `%in%`. With this operator, we can select rows based on a prespecified vector of values. This can be useful if there are specified values (e.g., specific `USUBJID`) which we would like to look at.

```r
# we save 4 USUBJID's in a vector:
lookup_ids <- c("01-716-1151", "01-710-1443", "01-708-1184", "01-705-1186")


# and then create a logical vector which returns TRUE for every entry in the
```

```r
# USUBJID vector which are represented in the lookup_ids, and else FALSE
adsl$USUBJID %in% lookup_ids
```

```
  [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
 [13] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
 [25] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
 [37] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
 [49] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
 [61] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
 [73] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
 [85] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
 [97] FALSE FALSE FALSE FALSE FALSE FALSE  TRUE FALSE FALSE FALSE FALSE FALSE
[109] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[121] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[133] FALSE FALSE FALSE FALSE  TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[145] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[157] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[169] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[181] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[193] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[205] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[217] FALSE  TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[229] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[241] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[253] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[265] FALSE FALSE FALSE FALSE  TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[277] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[289] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[301] FALSE FALSE FALSE FALSE FALSE FALSE
```

```r
# this approach can be used in the filter function:
adsl %>%
  select(STUDYID, USUBJID, ARM, AGE, SEX, RACE) %>%
  filter(USUBJID %in% lookup_ids)
```

```
# A tibble: 4 x 6
  STUDYID     USUBJID     ARM               AGE SEX   RACE
  <chr>       <chr>       <chr>           <dbl> <chr> <chr>
1 CDISCPILOT01 01-705-1186 Placebo           84 F     WHITE
2 CDISCPILOT01 01-708-1184 Screen Failure    70 F     WHITE
```

```
3 CDISCPILOT01 01-710-1443 Screen Failure      88 F      WHITE
4 CDISCPILOT01 01-716-1151 Xanomeline Low Dose  83 F      WHITE
```

Note that within the `filter` function (and in all major `dplyr` functions) R looks for the requested variables first within the supplied data-frame and afterwards in the global environment.

4.1.3 arrange

We can sort the dataframe with the arrange() function. It allows the sorting based on multiple variables. Note that the order of arranging variables determines the sorting hierarchy, so in this example we first order by `AGE` and

```
adsl %>%
  select(STUDYID, USUBJID, ARM, AGE, SEX, RACE) %>%
  filter(AGE >= 70,
         SEX == "F") %>%
  arrange(ARM, AGE)
```

```
# A tibble: 141 x 6
   STUDYID     USUBJID     ARM       AGE SEX   RACE
   <chr>       <chr>       <chr>   <dbl> <chr> <chr>
 1 CDISCPILOT01 01-705-1282 Placebo    70 F     BLACK OR AFRICAN AMERICAN
 2 CDISCPILOT01 01-704-1260 Placebo    71 F     WHITE
 3 CDISCPILOT01 01-703-1210 Placebo    72 F     WHITE
 4 CDISCPILOT01 01-716-1026 Placebo    73 F     WHITE
 5 CDISCPILOT01 01-718-1150 Placebo    73 F     WHITE
 6 CDISCPILOT01 01-708-1087 Placebo    74 F     WHITE
 7 CDISCPILOT01 01-708-1316 Placebo    74 F     WHITE
 8 CDISCPILOT01 01-709-1001 Placebo    76 F     WHITE
 9 CDISCPILOT01 01-710-1077 Placebo    76 F     WHITE
10 CDISCPILOT01 01-715-1397 Placebo    76 F     WHITE
# ... with 131 more rows
```

To sort by descending order, we can use the helper function `desc()` within `arrange()`:

```
adsl %>%
  select(STUDYID, USUBJID, ARM, AGE, SEX, RACE) %>%
  filter(AGE >= 70,
         SEX == "F") %>%
```

```
  arrange(ARM, desc(AGE))
```

```
# A tibble: 141 x 6
   STUDYID      USUBJID      ARM        AGE SEX   RACE
   <chr>        <chr>        <chr>    <dbl> <chr> <chr>
 1 CDISCPILOT01 01-710-1083 Placebo     89 F     WHITE
 2 CDISCPILOT01 01-710-1368 Placebo     88 F     WHITE
 3 CDISCPILOT01 01-714-1035 Placebo     88 F     WHITE
 4 CDISCPILOT01 01-701-1387 Placebo     87 F     WHITE
 5 CDISCPILOT01 01-704-1233 Placebo     87 F     WHITE
 6 CDISCPILOT01 01-716-1024 Placebo     87 F     WHITE
 7 CDISCPILOT01 01-705-1349 Placebo     86 F     WHITE
 8 CDISCPILOT01 01-710-1271 Placebo     86 F     WHITE
 9 CDISCPILOT01 01-716-1108 Placebo     86 F     WHITE
10 CDISCPILOT01 01-701-1047 Placebo     85 F     WHITE
# ... with 131 more rows
```

5 dplyr exercises

Stefan Thoma

```r
library("tidyverse")
```

```
-- Attaching packages --------------------------------------- tidyverse 1.3.2 --
v ggplot2 3.3.6      v purrr   0.3.5
v tibble  3.1.8      v dplyr   1.0.10
v tidyr   1.2.1      v stringr 1.4.1
v readr   2.1.3      v forcats 0.5.2
-- Conflicts ------------------------------------------ tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
```

```r
# load data
adsl <- read_csv("data/adsl.csv")
```

```
Rows: 306 Columns: 50
-- Column specification --------------------------------------------------------
Delimiter: ","
chr  (23): STUDYID, USUBJID, DTHFL, AGEU, SEX, RACE, ETHNIC, ARMCD, ARM, ACT...
dbl   (7): SUBJID, SITEID, AGE, DMDY, TRTDURD, DTHADY, LDDTHELD
lgl   (3): RFICDTC, REGION1, DTHA30FL
dttm  (3): RFPENDTC, TRTSDTM, TRTEDTM
date (14): RFSTDTC, RFENDTC, RFXSTDTC, RFXENDTC, DTHDTC, DMDTC, TRTSDT, TRTE...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

5.1 Data wrangling with dplyr

Load the `adsl` data-frame and select the following variables:

- USUBJID

- ARM

- SEX

- AGE

- AGEU

- AGEGR1

- COUNTRY

- EOSSTT

```
# we use starts_with("AGE") because we want to include every AGE variable that is in the o
adsl %>%
  select(USUBJID, ARM, SEX, starts_with("AGE"), COUNTRY, EOSSTT)
```

```
# A tibble: 306 x 8
   USUBJID     ARM                 SEX    AGE AGEU  AGEGR1 COUNTRY EOSSTT
   <chr>       <chr>               <chr> <dbl> <chr> <chr>  <chr>   <chr>
 1 01-701-1015 Placebo             F        63 YEARS 18-64  USA     COMPLETED
 2 01-701-1023 Placebo             M        64 YEARS 18-64  USA     DISCONTINU~
 3 01-701-1028 Xanomeline High Dose M       71 YEARS >=65   USA     COMPLETED
 4 01-701-1033 Xanomeline Low Dose  M       74 YEARS >=65   USA     DISCONTINU~
 5 01-701-1034 Xanomeline High Dose F       77 YEARS >=65   USA     COMPLETED
 6 01-701-1047 Placebo             F        85 YEARS >=65   USA     DISCONTINU~
 7 01-701-1057 Screen Failure      F        59 YEARS 18-64  USA     <NA>
 8 01-701-1097 Xanomeline Low Dose  M       68 YEARS >=65   USA     COMPLETED
 9 01-701-1111 Xanomeline Low Dose  F       81 YEARS >=65   USA     DISCONTINU~
10 01-701-1115 Xanomeline Low Dose  M       84 YEARS >=65   USA     DISCONTINU~
# ... with 296 more rows
```

On the selected variables, include only patients in the placebo arm who are 66, 77, 88, or 99 years old.

```
# There are different ways to solve this. The best way to filter the AGE is to create a ve
age_vec <- c(66, 77, 88, 99)
# or
```

```
  age_vec  <- 6:9 * 11
  # we can then use either the age_vec or the code that created the age_vec directly as a st


  adsl %>%
    select(USUBJID, SEX, ARM, EOSSTT, starts_with("AGE")) %>%
    filter(ARM == "Placebo",
           AGE %in% c(66, 77, 88, 99))
```

```
# A tibble: 5 x 7
  USUBJID       SEX   ARM     EOSSTT          AGE AGEU  AGEGR1
  <chr>         <chr> <chr>   <chr>         <dbl> <chr> <chr>
1 01-705-1059 F       Placebo DISCONTINUED     66 YEARS >=65
2 01-708-1171 F       Placebo COMPLETED        77 YEARS >=65
3 01-710-1368 F       Placebo COMPLETED        88 YEARS >=65
4 01-714-1035 F       Placebo COMPLETED        88 YEARS >=65
5 01-718-1139 M       Placebo COMPLETED        77 YEARS >=65
```

Further include the variable TRTSDTM (datetime of first exposure to treatment) and sort the previous data-frame according to this variable from most recent to least recent first exposure.

```
  adsl %>%
    select(USUBJID, SEX, ARM, EOSSTT, starts_with("AGE"), TRTSDTM) %>%
    filter(ARM == "Placebo",
           AGE %in% c(66, 77, 88, 99)) %>%
    arrange(desc(TRTSDTM))
```

```
# A tibble: 5 x 8
  USUBJID       SEX   ARM     EOSSTT          AGE AGEU  AGEGR1 TRTSDTM
  <chr>         <chr> <chr>   <chr>         <dbl> <chr> <chr>  <dttm>
1 01-714-1035 F       Placebo COMPLETED        88 YEARS >=65   2014-04-17 00:00:00
2 01-710-1368 F       Placebo COMPLETED        88 YEARS >=65   2013-10-23 00:00:00
3 01-705-1059 F       Placebo DISCONTINUED     66 YEARS >=65   2013-08-05 00:00:00
4 01-718-1139 M       Placebo COMPLETED        77 YEARS >=65   2013-05-19 00:00:00
5 01-708-1171 F       Placebo COMPLETED        77 YEARS >=65   2012-12-06 00:00:00
```

Part III

mutate

6 mutate

Creating New Columns Using `mutate()`

Thomas

```r
library(dplyr)
library(lubridate)
dm <- readRDS("data/dm.rds")
ae <- readRDS("data/ae.rds")
```

The equivalent of creating a new variables in SAS inside a `data` step is to use the `mutate()` function. In the simplest case a static value is assigned to the new column.

```r
adsl <- dm %>% mutate(DATASET = "ADSL")
```

This will set the value of the new variable `DATASET` to `"ADSL"` for all records.

```r
adsl %>% select(DATASET)
```

```
   DATASET
   <chr>
 1 ADSL
 2 ADSL
 3 ADSL
 4 ADSL
 5 ADSL
 6 ADSL
 7 ADSL
 8 ADSL
 9 ADSL
10 ADSL
# ... with 296 more rows
```

Note that new variables are always appended after existing columns such that `DATASET` is now the last column of `adsl`.

```
colnames(adsl)
```

```
 [1] "STUDYID"  "DOMAIN"   "USUBJID"  "SUBJID"   "RFSTDTC"  "RFENDTC"
 [7] "RFXSTDTC" "RFXENDTC" "RFICDTC"  "RFPENDTC" "DTHDTC"   "DTHFL"
[13] "SITEID"   "AGE"      "AGEU"     "SEX"      "RACE"     "ETHNIC"
[19] "ARMCD"    "ARM"      "ACTARMCD" "ACTARM"   "COUNTRY"  "DMDTC"
[25] "DMDY"     "DATASET"
```

Assigning the value of an existing column to a new column is the same as in SAS. The new column name goes to the left of `=` and the existing column to the right.

```
adsl <- adsl %>% mutate(TRT01P = ARM)
adsl %>% select(ARM, TRT01P)
```

```
# A tibble: 306 x 2
   ARM                 TRT01P
   <chr>               <chr>
 1 Placebo             Placebo
 2 Placebo             Placebo
 3 Xanomeline High Dose Xanomeline High Dose
 4 Xanomeline Low Dose  Xanomeline Low Dose
 5 Xanomeline High Dose Xanomeline High Dose
 6 Placebo             Placebo
 7 Screen Failure      Screen Failure
 8 Xanomeline Low Dose  Xanomeline Low Dose
 9 Xanomeline Low Dose  Xanomeline Low Dose
10 Xanomeline Low Dose  Xanomeline Low Dose
# ... with 296 more rows
```

In most cases new variables are created by applying functions on existing variables to somehow transform them.

```
adsl <- adsl %>% mutate(RFSTDT = ymd(RFSTDTC))
```

You can create multiple new variables inside `mutate()` similar to how you would do it inside a `data` step.

```
adae <- ae %>% mutate(
  ASTDT = ymd(AESTDTC),
  ASTDY = ASTDT - TRTSDT + 1
)
adae %>% select(AESTDTC, ASTDT, TRTSDT, ASTDY)
```

```
# A tibble: 1,191 x 4
   AESTDTC    ASTDT      TRTSDT     ASTDY
   <chr>      <date>     <date>     <drtn>
 1 2014-01-03 2014-01-03 2014-01-02  2 days
 2 2014-01-03 2014-01-03 2014-01-02  2 days
 3 2014-01-09 2014-01-09 2014-01-02  8 days
 4 2012-08-26 2012-08-26 2012-08-05 22 days
 5 2012-08-07 2012-08-07 2012-08-05  3 days
 6 2012-08-07 2012-08-07 2012-08-05  3 days
 7 2012-08-07 2012-08-07 2012-08-05  3 days
 8 2013-07-21 2013-07-21 2013-07-19  3 days
 9 2013-08-08 2013-08-08 2013-07-19 21 days
10 2014-08-27 2014-08-27 2014-07-01 58 days
# ... with 1,181 more rows
```

Just like in SAS you can use conditional logic to assign different values to a new variable depending on which value another variable has using `if_else()`.

```
adae %>%
  mutate(ASTDY = if_else(ASTDT <= TRTSDT, ASTDT - TRTSDT, ASTDT - TRTSDT + 1)) %>%
  select(USUBJID, TRTSDT, ASTDT, ASTDY)
```

```
# A tibble: 1,191 x 4
   USUBJID     TRTSDT     ASTDT      ASTDY
   <chr>       <date>     <date>     <drtn>
 1 01-701-1015 2014-01-02 2014-01-03  2 days
 2 01-701-1015 2014-01-02 2014-01-03  2 days
 3 01-701-1015 2014-01-02 2014-01-09  8 days
 4 01-701-1023 2012-08-05 2012-08-26 22 days
 5 01-701-1023 2012-08-05 2012-08-07  3 days
 6 01-701-1023 2012-08-05 2012-08-07  3 days
 7 01-701-1023 2012-08-05 2012-08-07  3 days
 8 01-701-1028 2013-07-19 2013-07-21  3 days
 9 01-701-1028 2013-07-19 2013-08-08 21 days
10 01-701-1034 2014-07-01 2014-08-27 58 days
```

```
# ... with 1,181 more rows
```

At this point let's make a small excursion to cover how R handles missing values, i.e. `NA`, when using conditional logic. Unlike in SAS where missing numbers are the smallest possible values such that `. < 10` is true, in R any comparison involving `NA` returns `NA` as a result.

```
NA < 9
```

```
[1] NA
```

```
NA == 0
```

```
[1] NA
```

This is the same when using `if_else()`.

```
adsl$AGE[1] <- NA
adsl %>%
  mutate(AGEGR = if_else(AGE >= 65, "Elderly", "Adult")) %>%
  select(USUBJID, AGE, AGEGR)
```

```
# A tibble: 306 x 3
   USUBJID        AGE AGEGR
   <chr>        <int> <chr>
 1 01-701-1015    NA <NA>
 2 01-701-1023    64 Adult
 3 01-701-1028    71 Elderly
 4 01-701-1033    74 Elderly
 5 01-701-1034    77 Elderly
 6 01-701-1047    85 Elderly
 7 01-701-1057    59 Adult
 8 01-701-1097    68 Elderly
 9 01-701-1111    81 Elderly
10 01-701-1115    84 Elderly
# ... with 296 more rows
```

To check whether a value is missing use the `is.na()` function.

```
is.na(NA)
```

```
[1] TRUE
```

```
  is.na("NA")
```

```
[1] FALSE
```

Finally, it's noteworthy that is are actually different types on `NA`s in R. We'll make use of them next.

Table 6.1: Types of 'NA' in R

Type	Example	Missing Value
character	"Brazil"	NA_character_
double	2.51	NA_real_
integer	1L	NA_integer_
logical	FALSE	NA

If the logic is more complex than a simple `if_else()` than use `case_when()` instead.

```r
  adsl %>%
    mutate(
      AGEGR1 = case_when(
        AGE < 18 ~ "<18",
        AGE < 45 ~ "<45",
        AGE < 65 ~ "<65",
        TRUE ~ ">=65"
      )
    ) %>%
    select(USUBJID, AGE, AGEGR1)
```

```
# A tibble: 306 x 3
   USUBJID         AGE AGEGR1
   <chr>         <int> <chr>
 1 01-701-1015      NA >=65
 2 01-701-1023      64 <65
 3 01-701-1028      71 >=65
 4 01-701-1033      74 >=65
 5 01-701-1034      77 >=65
 6 01-701-1047      85 >=65
```

```
 7 01-701-1057    59 <65
 8 01-701-1097    68 >=65
 9 01-701-1111    81 >=65
10 01-701-1115    84 >=65
# ... with 296 more rows
```

The final condition `TRUE` is the is a catch all term and must be used with some caution. Notice what happened to the `AGE` of the first subject whose value we set to `NA` above.

To mitigate this you should either explicitly handle missing values as a separate condition or be explicit for all cases. The former would look something like this.

```
adsl %>%
  mutate(
    AGEGR1 = case_when(
      is.na(AGE) ~ NA_character_,
      AGE < 18 ~ "<18",
      AGE < 45 ~ "<45",
      AGE < 65 ~ "<65",
      TRUE ~ ">=65"
    )
  ) %>%
  select(USUBJID, AGE, AGEGR1)
```

```
# A tibble: 306 x 3
   USUBJID        AGE AGEGR1
   <chr>        <int> <chr>
 1 01-701-1015     NA <NA>
 2 01-701-1023     64 <65
 3 01-701-1028     71 >=65
 4 01-701-1033     74 >=65
 5 01-701-1034     77 >=65
 6 01-701-1047     85 >=65
 7 01-701-1057     59 <65
 8 01-701-1097     68 >=65
 9 01-701-1111     81 >=65
10 01-701-1115     84 >=65
# ... with 296 more rows
```

And the latter like this.

```
adsl %>%
  mutate(
    AGEGR1 = case_when(
      AGE < 18 ~ "<18",
      AGE < 45 ~ "<45",
      AGE < 65 ~ "<65",
      AGE >= 65 ~ ">=65"
    )
  ) %>%
  select(USUBJID, AGE, AGEGR1)
```

```
# A tibble: 306 x 3
   USUBJID        AGE AGEGR1
   <chr>        <int> <chr>
 1 01-701-1015     NA <NA>
 2 01-701-1023     64 <65
 3 01-701-1028     71 >=65
 4 01-701-1033     74 >=65
 5 01-701-1034     77 >=65
 6 01-701-1047     85 >=65
 7 01-701-1057     59 <65
 8 01-701-1097     68 >=65
 9 01-701-1111     81 >=65
10 01-701-1115     84 >=65
# ... with 296 more rows
```

Finally, note that when a value does not match any of the conditions given which may be the case when not using a final TRUE then it is assigned NA.

```
adsl %>%
  mutate(
    AGEGR1 = case_when(
      AGE < 18 ~ "<18",
      AGE < 45 ~ "<45",
      AGE < 65 ~ "<65"
    )
  ) %>%
  select(USUBJID, AGE, AGEGR1)
```

```
# A tibble: 306 x 3
   USUBJID        AGE AGEGR1
```

```
      <chr>        <int> <chr>
 1 01-701-1015    NA <NA>
 2 01-701-1023    64 <65
 3 01-701-1028    71 <NA>
 4 01-701-1033    74 <NA>
 5 01-701-1034    77 <NA>
 6 01-701-1047    85 <NA>
 7 01-701-1057    59 <65
 8 01-701-1097    68 <NA>
 9 01-701-1111    81 <NA>
10 01-701-1115    84 <NA>
# ... with 296 more rows
```

7 mutate exercises

Thomas Neitmann

```
library(dplyr)
library(lubridate)
dm <- readRDS("data/dm.rds")
ae <- readRDS("data/ae.rds")
```

7.1 Exercise 1

A treatment emergent adverse event is defined as an adverse event whose start date is on or after the treatment start date (TRTSDT) and at the latest starts 7 days after the treatment end date (TRTEDT). Given this definition calculate TRTEMFL.

Hint: Turn the --DTC variables into proper dates first using the ymd() function.

```
ae %>%
  mutate(
    ASTDT = ymd(AESTDTC),
    AENDT = ymd(AEENDTC),
    TRTEMFL = if_else(ASTDT >= TRTSDT & ASTDT <= TRTEDT + 7, "Y", NA_character_)
  ) %>%
  select(USUBJID, ASTDT, AENDT, TRTSDT, TRTEDT, TRTEMFL)
```

Warning: 19 failed to parse.

```
# A tibble: 1,191 x 6
   USUBJID     ASTDT      AENDT    TRTSDT     TRTEDT     TRTEMFL
   <chr>       <date>     <date>   <date>     <date>     <chr>
 1 01-701-1015 2014-01-03 NA       2014-01-02 2014-07-02 Y
 2 01-701-1015 2014-01-03 NA       2014-01-02 2014-07-02 Y
```

```
 3 01-701-1015 2014-01-09 2014-01-11 2014-01-02 2014-07-02 Y
 4 01-701-1023 2012-08-26 NA         2012-08-05 2012-09-01 Y
 5 01-701-1023 2012-08-07 2012-08-30 2012-08-05 2012-09-01 Y
 6 01-701-1023 2012-08-07 NA         2012-08-05 2012-09-01 Y
 7 01-701-1023 2012-08-07 2012-08-30 2012-08-05 2012-09-01 Y
 8 01-701-1028 2013-07-21 NA         2013-07-19 2014-01-14 Y
 9 01-701-1028 2013-08-08 NA         2013-07-19 2014-01-14 Y
10 01-701-1034 2014-08-27 NA         2014-07-01 2014-12-30 Y
# ... with 1,181 more rows
```

7.2 Exercise 2

Create a new variable REGION1 based upon COUTRY as shown in the table below.

Countries	Region
Mexico, USA, Canada	North America
Spain, Greece, Germany, Switzerland	Europe
China, Japan	Asia

```
dm %>%
  mutate(
    REGION1 = case_when(
      COUNTRY %in% c("Mexico", "USA", "Canada") ~ "North America",
      COUNTRY %in% c("Spain", "Greece", "Germany", "Switzerland") ~ "Europe",
      COUNTRY %in% c("China", "Japan") ~ "Asia"
    )
  ) %>%
  select(USUBJID, COUNTRY, REGION1)
```

```
# A tibble: 306 x 3
  USUBJID     COUNTRY REGION1
  <chr>       <chr>   <chr>
 1 01-701-1015 USA     North America
 2 01-701-1023 USA     North America
 3 01-701-1028 USA     North America
 4 01-701-1033 USA     North America
 5 01-701-1034 USA     North America
 6 01-701-1047 USA     North America
 7 01-701-1057 USA     North America
```

```
 8 01-701-1097 USA      North America
 9 01-701-1111 USA      North America
10 01-701-1115 USA      North America
# ... with 296 more rows
```

Part IV

summarize

8 summarizing data

Thomas Neitmann

```r
library(dplyr)
dm <- readRDS("data/dm.rds")
ae <- readRDS("data/ae.rds")
```

While `mutate()` adds a new variable for all existing records to a dataset, `summarize()` aggregates one or more columns of a dataset thereby "collapsing" it. In the simplest case a single variable is aggregated using a summary function such as `mean()`.

```r
dm %>%
  summarize(avg_age = mean(AGE, na.rm = TRUE))
```

```
  avg_age
    <dbl>
1    75.1
```

Just like you can create multiple variables inside a single call to `mutate()` you can aggregate multiple variables (or the same variable with multiple summary functions) inside `summarize()`.

```r
dm %>%
  summarize(
    avg_age = mean(AGE, na.rm = TRUE),
    median_age = median(AGE, na.rm = TRUE)
  )
```

```
# A tibble: 1 x 2
  avg_age median_age
    <dbl>      <dbl>
1    75.1         77
```

So far we aggregated only numeric variables. Another useful aggregation is counting the number of records.

```
dm %>%
  summarize(N = n())
```

```
# A tibble: 1 x 1
      N
  <int>
1   306
```

This becomes quite powerful when combining `summarize()` with `group_by()`. This should look rather familiar to you if you every aggregated data using `proc sql`.

```
dm %>%
  group_by(COUNTRY) %>%
  summarize(n = n()) %>%
  ungroup()
```

```
# A tibble: 9 x 2
  COUNTRY          n
  <chr>        <int>
1 Canada          86
2 China           55
3 Germany         49
4 Greece          13
5 Japan           12
6 Mexico           5
7 Spain            6
8 Switzerland      9
9 USA             71
```

Note that it is best practice to `ungroup()` the dataset after you aggregated it. Failing to do so can lead to some rather bogus error when continuing to manipulate the aggregated dataset, e.g. using `mutate()`.

`group_by()` and `summarize()` can be used with numeric variables as well. In addition one can group by more than a single variable.

```
dm %>%
  group_by(ARM, COUNTRY) %>%
```

```
    summarize(avg_age = mean(AGE, na.rm = TRUE)) %>%
    ungroup()
```

```
# A tibble: 32 x 3
   ARM            COUNTRY      avg_age
   <chr>          <chr>          <dbl>
 1 Placebo        Canada          79.2
 2 Placebo        China           70.5
 3 Placebo        Germany         75.2
 4 Placebo        Greece          75.8
 5 Placebo        Japan           70
 6 Placebo        Mexico          65
 7 Placebo        Spain           83
 8 Placebo        Switzerland     69.3
 9 Placebo        USA             74.9
10 Screen Failure Canada          75.1
# ... with 22 more rows
```

9 summarizing data exercises

Thomas Neitmann

```r
library(dplyr)
dm <- readRDS("data/dm.rds")
ae <- readRDS("data/ae.rds")
```

9.1 Exercise 1

Count the number of *overall* adverse events per subject and sort the output such that the subject with the highest overall number of adverse events appears first.

```r
ae %>%
  group_by(USUBJID) %>%
  summarise(n_ae = n()) %>%
  arrange(desc(n_ae))
```

```
# A tibble: 225 x 2
   USUBJID        n_ae
   <chr>         <int>
 1 01-701-1302     23
 2 01-717-1004     19
 3 01-704-1266     16
 4 01-709-1029     16
 5 01-718-1427     16
 6 01-701-1192     15
 7 01-701-1275     15
 8 01-709-1309     15
 9 01-713-1179     15
10 01-711-1143     14
# ... with 215 more rows
```

9.2 Exercise 2

Count the overall number of *serious* adverse events per treatment arm (`ACTARM`).

```
ae %>%
  filter(AESER == "Y") %>%
  group_by(ACTARM) %>%
  summarise(n = n())
```

```
# A tibble: 2 x 2
  ACTARM                   n
  <chr>                <int>
1 Xanomeline High Dose     1
2 Xanomeline Low Dose      2
```

9.3 Exercise 3

Find the lowest and highest `AGE` per treatment arm.

```
dm %>%
  group_by(ARM) %>%
  summarise(youngest = min(AGE, na.rm = TRUE), oldest = max(AGE, na.rm = TRUE))
```

```
# A tibble: 4 x 3
  ARM                  youngest oldest
  <chr>                   <int>  <int>
1 Placebo                    52     89
2 Screen Failure             50     89
3 Xanomeline High Dose       56     88
4 Xanomeline Low Dose        51     88
```

Part V

tidyr

10 tidyr

Transposing data using `tidyr`and Joins

Zelos Zhu

10.0.1 Some Context

As we know, data can often be represented in several ways. Multiple observations of a variable can be organized by rows or by columns.

Table A.

ID	Pre	Post
x	1	2
y	3	4

Table B.

ID	Time	Value
x	Pre	1
x	Post	2
y	Pre	3
y	Post	4

When observations are spread along a row as multiple columns, we refer to the data as being in "wide" format (See Table A). When observations are spread along a column as multiple rows, we refer to the data as being in "long" format (See Table B). SDTM data for the most part generally adheres to the "long" structure, but as programmers we need to know how to work with both to suit our needs.

To get the desired shape of data, there are two useful functions from the `tidyr` package to make this transformation, aptly named: `pivot_longer()` and `pivot_wider()` . These can be seen as the R-equivalent of `proc transpose` in SAS.

10.0.2 Setup

```r
library(dplyr)
library(tidyr)
suppdm <- readRDS("data/suppdm.rds") %>%
  select(USUBJID, QNAM, QVAL)

head(suppdm, 10)
```

```
# A tibble: 10 x 3
   USUBJID      QNAM      QVAL
   <chr>        <chr>     <chr>
 1 01-701-1015 COMPLT16  Y
 2 01-701-1015 COMPLT24  Y
 3 01-701-1015 COMPLT8   Y
 4 01-701-1015 EFFICACY  Y
 5 01-701-1015 ITT       Y
 6 01-701-1015 SAFETY    Y
 7 01-701-1023 EFFICACY  Y
 8 01-701-1023 ITT       Y
 9 01-701-1023 SAFETY    Y
10 01-701-1028 COMPLT16  Y
```

As we see here, in our SUPPDM domain, the data is currently in the "long" format. If we wanted to transform the dataset such that each of the unique values of QNAM was their own column, we are looking to transpose the data from "long" to "wide". In this case, we use `pivot_wider()`.

```r
suppdm_wide <- suppdm %>%
  pivot_wider(
    names_from = "QNAM", # assign column names based on QNAM
    values_from = "QVAL" # retrieve values from QVAL
  )
suppdm_wide
```

```
# A tibble: 254 x 7
   USUBJID      COMPLT16 COMPLT24 COMPLT8 EFFICACY ITT   SAFETY
   <chr>        <chr>    <chr>    <chr>   <chr>    <chr> <chr>
 1 01-701-1015 Y        Y        Y       Y        Y     Y
 2 01-701-1023 <NA>     <NA>     <NA>    Y        Y     Y
 3 01-701-1028 Y        Y        Y       Y        Y     Y
```

```
 4 01-701-1033 <NA>       <NA>       <NA>       Y        Y        Y
 5 01-701-1034 Y          Y          Y          Y        Y        Y
 6 01-701-1047 <NA>       <NA>       <NA>       Y        Y        Y
 7 01-701-1097 Y          Y          Y          Y        Y        Y
 8 01-701-1111 <NA>       <NA>       <NA>       Y        Y        Y
 9 01-701-1115 <NA>       <NA>       Y          Y        Y        Y
10 01-701-1118 Y          Y          Y          Y        Y        Y
# ... with 244 more rows
```

Voila! This "wide" dataset may prove useful for joins (to be discussed later). But for now, let's pretend that this "wide" format is how our original data came to us in. If we wanted to take these respective flagging columns and turn them into a "long" format, we use `pivot_longer()`.

```
suppdm_long <- suppdm_wide %>%
  pivot_longer(
    cols = c("COMPLT16", "COMPLT24", "COMPLT8", "EFFICACY", "ITT", "SAFETY"),
    names_to = "QNAM",
    values_to = "QVAL"
  )
suppdm_long
```

```
# A tibble: 1,524 x 3
   USUBJID      QNAM      QVAL
   <chr>        <chr>     <chr>
 1 01-701-1015 COMPLT16 Y
 2 01-701-1015 COMPLT24 Y
 3 01-701-1015 COMPLT8  Y
 4 01-701-1015 EFFICACY Y
 5 01-701-1015 ITT      Y
 6 01-701-1015 SAFETY   Y
 7 01-701-1023 COMPLT16 <NA>
 8 01-701-1023 COMPLT24 <NA>
 9 01-701-1023 COMPLT8  <NA>
10 01-701-1023 EFFICACY Y
# ... with 1,514 more rows
```

As you can see, as we pivoted back, we didn't come up with an *exact* duplicate of our original `suppdm` dataframe. This is because the default of `pivot_longer()` is **not** to drop NA values, which can be modified with the `values_drop_na` function input, just one of the many powerful additional function inputs from both of these pivoting functions. `pivot_wider()` and

`pivot_longer()` were designed to handle a variety of situations when transposing data in the most flexible of ways.

```
suppdm_long <- suppdm_wide %>%
  pivot_longer(
    cols = c("COMPLT16", "COMPLT24", "COMPLT8", "EFFICACY", "ITT", "SAFETY"),
    names_to = "flag",
    values_to = "flag_value",
    values_drop_na = TRUE
  )
suppdm_long
```

```
# A tibble: 1,197 x 3
   USUBJID      flag      flag_value
   <chr>        <chr>     <chr>
 1 01-701-1015 COMPLT16 Y
 2 01-701-1015 COMPLT24 Y
 3 01-701-1015 COMPLT8  Y
 4 01-701-1015 EFFICACY Y
 5 01-701-1015 ITT      Y
 6 01-701-1015 SAFETY   Y
 7 01-701-1023 EFFICACY Y
 8 01-701-1023 ITT      Y
 9 01-701-1023 SAFETY   Y
10 01-701-1028 COMPLT16 Y
# ... with 1,187 more rows
```

Bonus Trick: The `names_to`/`values_to` function arguments can prove to be helpful as a renaming step during the data cleaning process too!

10.1 Relational Data (Joins)

When a pair of tables need to be joined together, we have a variety of functions that can achieve such a task:

- `left_join()`

- `right_join()`

- `full_join()`

- `inner_join()`

The use of these functions is very similar to `proc sql` in SAS. `left_join()` will cover most of use cases and is demonstrated below:

```r
dm <- admiral.test::admiral_dm %>%
  select(STUDYID, USUBJID, AGE, ARM)

dm_suppdm <- dm %>%
  left_join(suppdm_wide, by = "USUBJID")

head(dm_suppdm)
```

```
# A tibble: 6 x 10
  STUDYID     USUBJID     AGE ARM   COMPL~1 COMPL~2 COMPLT8 EFFIC~3 ITT   SAFETY
  <chr>       <chr>     <dbl> <chr> <chr>   <chr>   <chr>   <chr>   <chr> <chr>
1 CDISCPILOT01 01-701-~    63 Plac~ Y       Y       Y       Y       Y     Y
2 CDISCPILOT01 01-701-~    64 Plac~ <NA>    <NA>    <NA>    Y       Y     Y
3 CDISCPILOT01 01-701-~    71 Xano~ Y       Y       Y       Y       Y     Y
4 CDISCPILOT01 01-701-~    74 Xano~ <NA>    <NA>    <NA>    Y       Y     Y
5 CDISCPILOT01 01-701-~    77 Xano~ Y       Y       Y       Y       Y     Y
6 CDISCPILOT01 01-701-~    85 Plac~ <NA>    <NA>    <NA>    Y       Y     Y
# ... with abbreviated variable names 1: COMPLT16, 2: COMPLT24, 3: EFFICACY
```

The join can also be completed with different column names as long as you define the join-key relationship, demonstrated below:

```r
dummy1 <- data.frame(
  STUDYID = c("TRIALX", "TRIALX"),
  USUBJID = c("1001", "1002"),
  AGE = c(18, 22)
)

dummy2 <- data.frame(
  STUDYID = c("TRIALX", "TRIALX"),
  SUBJECT = c("1001", "1002"),
  SEX = c("M", "F")
)

dummy3 <- dummy1 %>%
  left_join(dummy2, by = c("STUDYID" = "STUDYID", "USUBJID" = "SUBJECT"))

head(dummy3)
```

```
  STUDYID USUBJID AGE SEX
1 TRIALX    1001  18   M
2 TRIALX    1002  22   F
```

11 tidyr exercises

Zelos Zhu

```
library(tidyverse)
library(dplyr)
library(tidyr)

# load data
ex <- readRDS("data/ex.rds")
dm <- readRDS("data/dm.rds")
ds <- readRDS("data/ds.rds")
suppds <- readRDS("data/suppds.rds")
```

11.1 Pivoting with tidyr

Load the `ex` data-frame from `admiral_ex`and select the following variables:

- USUBJID

- EXTRT

- VISIT

- EXSTDTC

```
ex %>%
  select(USUBJID, EXTRT, VISIT, EXSTDTC)
```

```
# A tibble: 591 x 4
   USUBJID     EXTRT      VISIT     EXSTDTC
   <chr>       <chr>      <chr>     <chr>
 1 01-701-1015 PLACEBO    BASELINE  2014-01-02
```

```
 2 01-701-1015 PLACEBO     WEEK 2   2014-01-17
 3 01-701-1015 PLACEBO     WEEK 24  2014-06-19
 4 01-701-1023 PLACEBO     BASELINE 2012-08-05
 5 01-701-1023 PLACEBO     WEEK 2   2012-08-28
 6 01-701-1028 XANOMELINE BASELINE 2013-07-19
 7 01-701-1028 XANOMELINE WEEK 2   2013-08-02
 8 01-701-1028 XANOMELINE WEEK 24  2014-01-07
 9 01-701-1033 XANOMELINE BASELINE 2014-03-18
10 01-701-1034 XANOMELINE BASELINE 2014-07-01
# ... with 581 more rows
```

Using `pivot_wider()` create a table that would shaped this way

USUBJID	EXTRT	BASELINE	WEEK 2	WEEK 24
...	...	...	...	...

```
  ex %>%
    select(USUBJID, EXTRT, VISIT, EXSTDTC) %>%
    pivot_wider(names_from = "VISIT", values_from = "EXSTDTC")
```

```
# A tibble: 254 x 5
   USUBJID     EXTRT      BASELINE   `WEEK 2`   `WEEK 24`
   <chr>       <chr>      <chr>      <chr>      <chr>
 1 01-701-1015 PLACEBO    2014-01-02 2014-01-17 2014-06-19
 2 01-701-1023 PLACEBO    2012-08-05 2012-08-28 <NA>
 3 01-701-1028 XANOMELINE 2013-07-19 2013-08-02 2014-01-07
 4 01-701-1033 XANOMELINE 2014-03-18 <NA>       <NA>
 5 01-701-1034 XANOMELINE 2014-07-01 2014-07-16 2014-12-18
 6 01-701-1047 PLACEBO    2013-02-12 2013-02-26 <NA>
 7 01-701-1097 XANOMELINE 2014-01-01 2014-01-16 2014-06-19
 8 01-701-1111 XANOMELINE 2012-09-07 <NA>       <NA>
 9 01-701-1115 XANOMELINE 2012-11-30 2012-12-14 <NA>
10 01-701-1118 PLACEBO    2014-03-12 2014-03-27 2014-08-28
# ... with 244 more rows
```

Load the `dm` data-frame from `admiral_dm` and select the following variables:

- USUBJID

- RACE

- SEX

```
  dm %>%
    select(USUBJID, RACE, SEX)
```

```
# A tibble: 306 x 3
   USUBJID      RACE  SEX
   <chr>        <chr> <chr>
 1 01-701-1015 WHITE F
 2 01-701-1023 WHITE M
 3 01-701-1028 WHITE M
 4 01-701-1033 WHITE M
 5 01-701-1034 WHITE F
 6 01-701-1047 WHITE F
 7 01-701-1057 WHITE F
 8 01-701-1097 WHITE M
 9 01-701-1111 WHITE F
10 01-701-1115 WHITE M
# ... with 296 more rows
```

Using `pivot_longer()` create a table that would shaped this way

USUBJID	VAR	VAL
1001	RACE	WHITE
1001	SEX	M

```
  dm %>%
    select(USUBJID, RACE, SEX) %>%
    pivot_longer(cols = c(RACE, SEX),
                 names_to = "VAR",
                 values_to = "VAL")
```

```
# A tibble: 612 x 3
   USUBJID      VAR   VAL
   <chr>        <chr> <chr>
 1 01-701-1015 RACE  WHITE
 2 01-701-1015 SEX   F
 3 01-701-1023 RACE  WHITE
 4 01-701-1023 SEX   M
 5 01-701-1028 RACE  WHITE
 6 01-701-1028 SEX   M
```

```
 7 01-701-1033 RACE  WHITE
 8 01-701-1033 SEX   M
 9 01-701-1034 RACE  WHITE
10 01-701-1034 SEX   F
# ... with 602 more rows
```

11.2 Joining using dplyr

Load the `ds` data-frame from `admiral_ds` and `suppds` data-frame from `admiral_suppds`. Prior to joining the two datasets together, we may need to do some cleaning of the data on `suppds`.

- Filter IDVAR for "DSSEQ"

- Mutate IDVARVAL from type character to type numeric.

- Select USUBJID IDVARVAL QNAM QLABEL QVAL

```r
suppds <- suppds %>%
  filter(IDVAR == "DSSEQ") %>%
  mutate(IDVARVAL = as.numeric(IDVARVAL)) %>%
  select(USUBJID, IDVARVAL, QNAM, QLABEL, QVAL)

suppds
```

```
# A tibble: 3 x 5
  USUBJID      IDVARVAL QNAM    QLABEL                           QVAL
  <chr>           <dbl> <chr>   <chr>                            <chr>
1 01-703-1175         2 ENTCRIT PROTOCOL ENTRY CRITERIA NOT MET 16
2 01-705-1382         2 ENTCRIT PROTOCOL ENTRY CRITERIA NOT MET 25
3 01-708-1372         3 ENTCRIT PROTOCOL ENTRY CRITERIA NOT MET 16
```

Join the two tables together using `USUBJID` and `DSSEQ` as the key joining variables.

```r
ds %>%
  left_join(suppds, by = c("USUBJID" = "USUBJID", "DSSEQ" = "IDVARVAL"))
```

```
# A tibble: 850 x 16
   STUDYID   DOMAIN USUBJID DSSEQ DSSPID DSTERM DSDECOD DSCAT VISIT~1 VISIT DSDTC
   <chr>     <chr>  <chr>   <dbl> <chr>  <chr>  <chr>   <chr>   <dbl> <chr> <chr>
 1 CDISCPI~ DS     01-701~     1 <NA>   RANDO~ RANDOM~ PROT~       3 BASE~ 2014~
```

```
 2 CDISCPI~ DS      01-701~      2 <NA>    PROTO~ COMPLE~ DISP~       13 WEEK~ 2014~
 3 CDISCPI~ DS      01-701~      3 <NA>    FINAL~ FINAL ~ OTHE~       13 WEEK~ 2014~
 4 CDISCPI~ DS      01-701~      1 <NA>    RANDO~ RANDOM~ PROT~        3 BASE~ 2012~
 5 CDISCPI~ DS      01-701~      2 24      ADVER~ ADVERS~ DISP~        5 WEEK~ 2012~
 6 CDISCPI~ DS      01-701~      3 <NA>    FINAL~ FINAL ~ OTHE~        5 WEEK~ 2012~
 7 CDISCPI~ DS      01-701~      4 <NA>    FINAL~ FINAL ~ OTHE~      201 RETR~ 2013~
 8 CDISCPI~ DS      01-701~      1 <NA>    RANDO~ RANDOM~ PROT~        3 BASE~ 2013~
 9 CDISCPI~ DS      01-701~      2 <NA>    PROTO~ COMPLE~ DISP~       13 WEEK~ 2014~
10 CDISCPI~ DS      01-701~      3 <NA>    FINAL~ FINAL ~ OTHE~       13 WEEK~ 2014~
# ... with 840 more rows, 5 more variables: DSSTDTC <chr>, DSSTDY <dbl>,
#   QNAM <chr>, QLABEL <chr>, QVAL <chr>, and abbreviated variable name
#   1: VISITNUM
```